

ex16

January 8, 2024

```
[22]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings("ignore")
# sns.set_style("darkgrid", {"grid.color": ".6",
#                             "grid.linestyle": ":"})

from sklearn.preprocessing import StandardScaler , MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
import seaborn as sns
from sklearn.metrics import accuracy_score
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn.tree import DecisionTreeClassifier

# from xgboost import XGBClassifier
# from sklearn.metrics import r2_score
# from sklearn.metrics import mean_squared_error
# from sklearn.model_selection import GridSearchCV
```

```
[2]: df = pd.read_csv('/Users/thutranghoa/Code/Data_analysis/Data/Breast Cancer_
↳Wisconsin (Diagnostic) Data Set.csv')
df
```

```
[2]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	
..	
564	926424	M	21.56	22.39	142.00	1479.0	

565	926682	M	20.13	28.25	131.20	1261.0
566	926954	M	16.60	28.08	108.30	858.1
567	927241	M	20.60	29.33	140.10	1265.0
568	92751	B	7.76	24.54	47.92	181.0

	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	\
0	0.11840	0.27760	0.30010	0.14710	
1	0.08474	0.07864	0.08690	0.07017	
2	0.10960	0.15990	0.19740	0.12790	
3	0.14250	0.28390	0.24140	0.10520	
4	0.10030	0.13280	0.19800	0.10430	
..	
564	0.11100	0.11590	0.24390	0.13890	
565	0.09780	0.10340	0.14400	0.09791	
566	0.08455	0.10230	0.09251	0.05302	
567	0.11780	0.27700	0.35140	0.15200	
568	0.05263	0.04362	0.00000	0.00000	

	texture_worst	perimeter_worst	area_worst	smoothness_worst	\
0	17.33	184.60	2019.0	0.16220	
1	23.41	158.80	1956.0	0.12380	
2	25.53	152.50	1709.0	0.14440	
3	26.50	98.87	567.7	0.20980	
4	16.67	152.20	1575.0	0.13740	
..	
564	26.40	166.10	2027.0	0.14100	
565	38.25	155.00	1731.0	0.11660	
566	34.12	126.70	1124.0	0.11390	
567	39.42	184.60	1821.0	0.16500	
568	30.37	59.16	268.6	0.08996	

	compactness_worst	concavity_worst	concave points_worst	symmetry_worst	\
0	0.66560	0.7119	0.2654	0.4601	
1	0.18660	0.2416	0.1860	0.2750	
2	0.42450	0.4504	0.2430	0.3613	
3	0.86630	0.6869	0.2575	0.6638	
4	0.20500	0.4000	0.1625	0.2364	
..	
564	0.21130	0.4107	0.2216	0.2060	
565	0.19220	0.3215	0.1628	0.2572	
566	0.30940	0.3403	0.1418	0.2218	
567	0.86810	0.9387	0.2650	0.4087	
568	0.06444	0.0000	0.0000	0.2871	

	fractal_dimension_worst	Unnamed: 32
0	0.11890	NaN
1	0.08902	NaN

```

2          0.08758          NaN
3          0.17300          NaN
4          0.07678          NaN
..          ...          ...
564        0.07115          NaN
565        0.06637          NaN
566        0.07820          NaN
567        0.12400          NaN
568        0.07039          NaN

```

[569 rows x 33 columns]

```
[12]: df = df.drop(['Unnamed: 32'], axis=1)
```

```
[18]: df['diagnosis'] = df['diagnosis'].map({'M': 1, 'B': 0})
```

```
[15]: df.shape
```

```
[15]: (569, 32)
```

```
[19]: df.describe()
```

```
[19]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	\
count	5.690000e+02	569.000000	569.000000	569.000000	569.000000	
mean	3.037183e+07	0.372583	14.127292	19.289649	91.969033	
std	1.250206e+08	0.483918	3.524049	4.301036	24.298981	
min	8.670000e+03	0.000000	6.981000	9.710000	43.790000	
25%	8.692180e+05	0.000000	11.700000	16.170000	75.170000	
50%	9.060240e+05	0.000000	13.370000	18.840000	86.240000	
75%	8.813129e+06	1.000000	15.780000	21.800000	104.100000	
max	9.113205e+08	1.000000	28.110000	39.280000	188.500000	

	area_mean	smoothness_mean	compactness_mean	concavity_mean	\
count	569.000000	569.000000	569.000000	569.000000	
mean	654.889104	0.096360	0.104341	0.088799	
std	351.914129	0.014064	0.052813	0.079720	
min	143.500000	0.052630	0.019380	0.000000	
25%	420.300000	0.086370	0.064920	0.029560	
50%	551.100000	0.095870	0.092630	0.061540	
75%	782.700000	0.105300	0.130400	0.130700	
max	2501.000000	0.163400	0.345400	0.426800	

	concave	points_mean	...	radius_worst	texture_worst	perimeter_worst	\
count		569.000000	...	569.000000	569.000000	569.000000	
mean		0.048919	...	16.269190	25.677223	107.261213	
std		0.038803	...	4.833242	6.146258	33.602542	
min		0.000000	...	7.930000	12.020000	50.410000	

25%	0.020310	...	13.010000	21.080000	84.110000
50%	0.033500	...	14.970000	25.410000	97.660000
75%	0.074000	...	18.790000	29.720000	125.400000
max	0.201200	...	36.040000	49.540000	251.200000

	area_worst	smoothness_worst	compactness_worst	concavity_worst	\
count	569.000000	569.000000	569.000000	569.000000	
mean	880.583128	0.132369	0.254265	0.272188	
std	569.356993	0.022832	0.157336	0.208624	
min	185.200000	0.071170	0.027290	0.000000	
25%	515.300000	0.116600	0.147200	0.114500	
50%	686.500000	0.131300	0.211900	0.226700	
75%	1084.000000	0.146000	0.339100	0.382900	
max	4254.000000	0.222600	1.058000	1.252000	

	concave points_worst	symmetry_worst	fractal_dimension_worst
count	569.000000	569.000000	569.000000
mean	0.114606	0.290076	0.083946
std	0.065732	0.061867	0.018061
min	0.000000	0.156500	0.055040
25%	0.064930	0.250400	0.071460
50%	0.099930	0.282200	0.080040
75%	0.161400	0.317900	0.092080
max	0.291000	0.663800	0.207500

[8 rows x 32 columns]

[20]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     569 non-null    int64
1   diagnosis                             569 non-null    int64
2   radius_mean                           569 non-null    float64
3   texture_mean                           569 non-null    float64
4   perimeter_mean                         569 non-null    float64
5   area_mean                             569 non-null    float64
6   smoothness_mean                        569 non-null    float64
7   compactness_mean                       569 non-null    float64
8   concavity_mean                         569 non-null    float64
9   concave points_mean                   569 non-null    float64
10  symmetry_mean                          569 non-null    float64
11  fractal_dimension_mean                 569 non-null    float64
12  radius_se                              569 non-null    float64
13  texture_se                             569 non-null    float64
```

```

14  perimeter_se          569 non-null    float64
15  area_se              569 non-null    float64
16  smoothness_se       569 non-null    float64
17  compactness_se      569 non-null    float64
18  concavity_se        569 non-null    float64
19  concave points_se    569 non-null    float64
20  symmetry_se         569 non-null    float64
21  fractal_dimension_se 569 non-null    float64
22  radius_worst        569 non-null    float64
23  texture_worst       569 non-null    float64
24  perimeter_worst     569 non-null    float64
25  area_worst          569 non-null    float64
26  smoothness_worst    569 non-null    float64
27  compactness_worst   569 non-null    float64
28  concavity_worst     569 non-null    float64
29  concave points_worst 569 non-null    float64
30  symmetry_worst      569 non-null    float64
31  fractal_dimension_worst 569 non-null    float64
dtypes: float64(30), int64(2)
memory usage: 142.4 KB

```

```
[43]: df.columns
```

```

[43]: Index(['id', 'diagnosis', 'radius_mean', 'texture_mean', 'perimeter_mean',
        'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean',
        'concave points_mean', 'symmetry_mean', 'fractal_dimension_mean',
        'radius_se', 'texture_se', 'perimeter_se', 'area_se', 'smoothness_se',
        'compactness_se', 'concavity_se', 'concave points_se', 'symmetry_se',
        'fractal_dimension_se', 'radius_worst', 'texture_worst',
        'perimeter_worst', 'area_worst', 'smoothness_worst',
        'compactness_worst', 'concavity_worst', 'concave points_worst',
        'symmetry_worst', 'fractal_dimension_worst'],
        dtype='object')

```

0.1 What are they ?

This dataset includes features about diagnosis of breast cancer. For details :

- ID : Patient identifier
- Diagnosis : indicates whether the tumor is Benign (B) or malignant (M) -> Output

** Data about physical dimensions of tumor (each is calculated 3 values : Mean, Worst and SE) -

Radius : the size of the tumor

- Texture : a measures of the tumor's surface smoothness

- Perimeter : linked to the size and growth rate of the tumor

- Area : Larger tumor areas are seen in advanced stages of breast cancer

** Data about tumor texture and composition (each is calculated 3 values : Mean, Worst and SE)

- Smoothness : variations in the tumor surface
- Compactness : correlate with denser, potentially cancerous
- Concavity : sign of malignancy
- Concavity points : increased number of concave points -> malignant tumor

** Others - Symmetry : asymmetrical tumors are more likely to be cancerous.


- Fractal dimension : higher values indicate more complex, malignant

0.2 Why are they chosen ?

- Because these features provide a detailed profile of the tumor, for accurate diagnosis and treatment planning
- Features like radius, perimeter and area are related to physical size of tumor, important for staging the cancer
- Texture, smoothness and compactness offer insights into tumor's cell structure, allow determining the type of cancer (B or M)

In conclusion, having and understanding features are essential for effective breast cancer management.

```
[21]: ' Box plot '
fig = plt.figure(figsize=(60, 50))
temp = df.columns.tolist()
for i, item in enumerate(temp):
    plt.subplot(8, 4, i+1)
    sns.boxplot(data=df, x=item, color='green')
plt.tight_layout(pad=0.4, w_pad=0.5, h_pad=2.0)
plt.show()
```



```
[23]: scaler = MinMaxScaler()
columns = df.columns
df = pd.DataFrame(scaler.fit_transform(df))
df.columns = columns
df.head()
```

```
[23]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	0.000915	1.0	0.521037	0.022658	0.545989	0.363733	
1	0.000915	1.0	0.643144	0.272574	0.615783	0.501591	
2	0.092495	1.0	0.601496	0.390260	0.595743	0.449417	
3	0.092547	1.0	0.210090	0.360839	0.233501	0.102906	
4	0.092559	1.0	0.629893	0.156578	0.630986	0.489290	

	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	\
0	0.593753	0.792037	0.703140	0.731113	
1	0.289880	0.181768	0.203608	0.348757	
2	0.514309	0.431017	0.462512	0.635686	
3	0.811321	0.811361	0.565604	0.522863	

4		0.430351	0.347893	0.463918	0.518390
---	--	----------	----------	----------	----------

	...	radius_worst	texture_worst	perimeter_worst	area_worst	\
0	...	0.620776	0.141525	0.668310	0.450698	
1	...	0.606901	0.303571	0.539818	0.435214	
2	...	0.556386	0.360075	0.508442	0.374508	
3	...	0.248310	0.385928	0.241347	0.094008	
4	...	0.519744	0.123934	0.506948	0.341575	

		smoothness_worst	compactness_worst	concavity_worst	concave points_worst	\
0		0.601136	0.619292	0.568610	0.912027	
1		0.347553	0.154563	0.192971	0.639175	
2		0.483590	0.385375	0.359744	0.835052	
3		0.915472	0.814012	0.548642	0.884880	
4		0.437364	0.172415	0.319489	0.558419	

		symmetry_worst	fractal_dimension_worst
0		0.598462	0.418864
1		0.233590	0.222878
2		0.403706	0.213433
3		1.000000	0.773711
4		0.157500	0.142595

[5 rows x 32 columns]


```
[26]: ' Heat map'
corr_matrix = df.corr()

# Set up the matplotlib figure
plt.figure(figsize=(20, 16))

# Draw the heatmap
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap='coolwarm')

# Add title
plt.title('Correlation Matrix Heatmap')

# Show the plot
plt.show()
```

```
[27]: X = df.drop('diagnosis', axis=1)
      y = df['diagnosis']
```

```
[28]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=42, stratify = y)
```

```
[29]: 'Model '

rf=RandomForestClassifier()
dt=DecisionTreeClassifier()
lr=LogisticRegression()
```

```
[30]: 'RF '
print ('Random Forest model : ')
rf.fit(X_train,y_train)
y_pred_rf=rf.predict(X_test)
```


```
print ('Train : ', accuracy_score(y_train,rf.predict(X_train)))
print ('Test : ', accuracy_score(y_test,y_pred_rf))
```

Random Forest model :
Train : 1.0
Test : 0.9649122807017544

```
[31]: from sklearn.metrics import ConfusionMatrixDisplay,confusion_matrix,␣
      ↪classification_report

disp = ConfusionMatrixDisplay.from_predictions(y_test, y_pred_rf)

plt.show()
```




```
[32]: 'Logistic '
print ('Logistic regression model : ')
lr.fit(X_train,y_train)
y_pred_lr=lr.predict(X_test)
print ('Train : ', accuracy_score(y_train,lr.predict(X_train)))
print ('Test : ', accuracy_score(y_test,y_pred_lr))
```

Logistic regression model :

Train : 0.9692307692307692
Test : 0.9736842105263158

```
[33]: disp = ConfusionMatrixDisplay.from_predictions(y_test, y_pred_lr)

plt.show()
```




```
[34]: 'DT '
print ('Decision Tree model : ')
dt.fit(X_train,y_train)
y_pred_dt=dt.predict(X_test)
print ('Train : ', accuracy_score(y_train,dt.predict(X_train)))
print ('Test : ', accuracy_score(y_test,y_pred_dt))
```

Decision Tree model :
Train : 1.0
Test : 0.9298245614035088

```
[35]: disp = ConfusionMatrixDisplay.from_predictions(y_test, y_pred_dt)

plt.show()
```



```
[36]: para={'n_estimators': [20, 60, 100, 120],
          'max_features' : [0.2, 0.6, 1.0],
          'max_depth' : [2, 8, None],
          'max_samples' : [0.5, 0.75, 1.0]
        }

# param_grid = {
#     'n_estimators': [200, 500],
#     'max_features': ['auto', 'sqrt', 'log2'],
#     'max_depth' : [4, 5, 6, 7, 8],
#     'criterion' : ['gini', 'entropy']
# }

rf = RandomForestClassifier()
CV_rfc = GridSearchCV(estimator=rf, param_grid=para, cv= 5)
CV_rfc.fit(X_train, y_train)

# rf_grid.fit(X_train, y_train)
```

```
[36]: GridSearchCV(cv=5, estimator=RandomForestClassifier(),
                  param_grid={'max_depth': [2, 8, None],
```

```
'max_features': [0.2, 0.6, 1.0],
'max_samples': [0.5, 0.75, 1.0],
'n_estimators': [20, 60, 100, 120]}}
```

```
[37]: print (CV_rfc.best_score_)
      'be improved not much'
```

0.9670329670329672

```
[42]: X = ['RF', 'DT', 'LR']
      y = [accuracy_score(y_test,y_pred_rf),
           ↪accuracy_score(y_test,y_pred_lr),accuracy_score(y_test,y_pred_dt)]

      plt.barh (X, y)

      for index, value in enumerate(y):
          plt.text(value, index,
                   str(value))
      plt.title('Accuracy of each model')
```

```
[42]: Text(0.5, 1.0, 'Accuracy of each model')
```

